

Kernel Programming Model



- Defines how the concurrency model is mapped to physical hardware.

```
void combine_two_arrays_CPU(float *A, float *B, float *C, int n) {  
    for (int i = 0; i < n; i++) {  
        C[i] = 1.0f / (sin(A[i])*cos(B[i]) + cos(A[i])*sin(B[i]));  
    }  
}
```




```
__kernel void CombineTwoArrays(__global float* A, __global float* B,  
                               __global float* C) {  
    int i = get_global_id(0);  
  
    C[i] = 1.0f / (sin(A[i])*cos(B[i]) + cos(A[i])*sin(B[i]));  
}
```

Instructions

```
Set i;  
R10 := x[i];  
R11 := y[i];  
R12 := sin(R10);  
R13 := cos(R11);  
R14 := R12*R13  
R12 := cos(R10);  
R13 := sin(R11);  
R14 := R12*R13 + R14;  
z[i] := 1.0/R14;
```

Execute in lock-step.





- Dimension, work-item, and work-group
 - **Work-item:** the unit of concurrent execution in OpenCL C
 - Work-items of n **dimension** are divided into smaller, equally sized, n -dimensional **work-groups**.
- Notes
 - Each work-group is assigned to a compute unit of device.
 - Each work-group can be assigned to the compute unit only when the device can afford the resources that the work-group needs.
 - **AMD GCN GPU**
 - Each work-group is divided into wavefronts of 64 work-items each.
 - Each wavefront is assigned to a SIMD unit within the compute unit.
 - Vector instructions executed by a wavefront are issued to the SIMD unit over four cycles.
 - Every four cycles, a new instruction can be issued to the SIMD unit.
 - Work-groups are processed in parallel in an unpredictable order.
 - Wavefronts in a work-group are processed in parallel in an unpredictable order.
 - Synchronization between work-items possible only within work-groups: barriers and memory fences.



- **1D example**

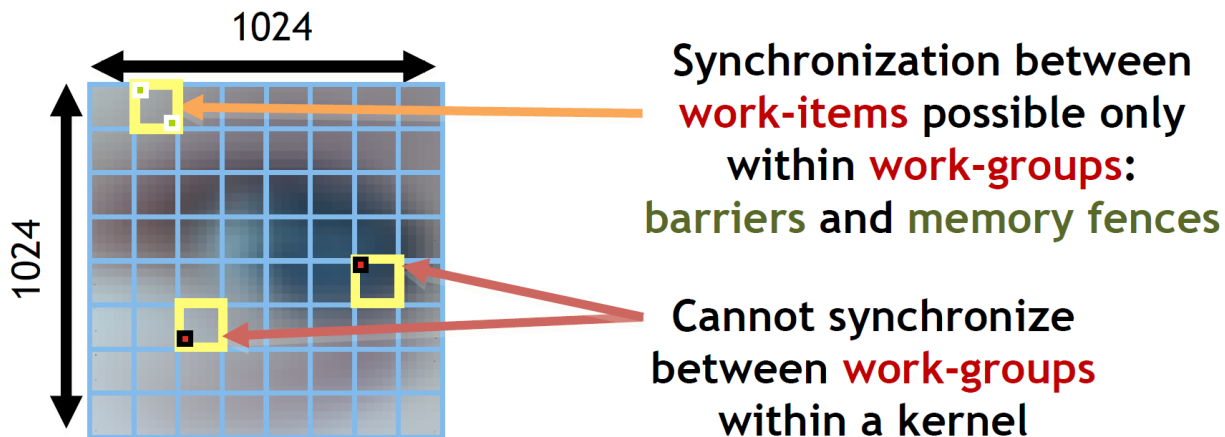
- 134,217,728 work-items/ 256 work-items per work-group

Work-items



- **2D example**

- 1,024x1,024 work-items/ 128x28 work-items per work-group





- How do a work-item get information about itself during kernel execution?

Table 5.2 Functions related to work-items

Function	Purpose
<code>uint get_work_dim()</code>	Returns the number of dimensions in the kernel's index space
<code>size_t get_global_size(uint dim)</code>	Returns the number of work-items for a given dimension
<code>size_t get_global_id(uint dim)</code>	Returns the element of the work-item's global ID for a given dimension
<code>size_t get_global_offset(uint dim)</code>	Returns the initial offset used to compute global IDs

Table 5.3 Functions related to work-groups

Function	Purpose
<code>size_t get_num_groups(uint dim)</code>	Returns the number of work-groups for a given dimension
<code>size_t get_group_id(uint dim)</code>	Returns the ID of the work-item's work-group for a given dimension
<code>size_t get_local_id(uint dim)</code>	Returns the ID of the work-item within its work-group for a given dimension
<code>size_t get_local_size(uint dim)</code>	Returns the number of work-items in the work-group for a given dimension

Example: Reduction

- 1,048,576 work-items in total
- 256 work-items per work-group
- 4,096 work-groups

Imagine how this OpenCL kernel runs over the work-items/work-groups/wavefronts.

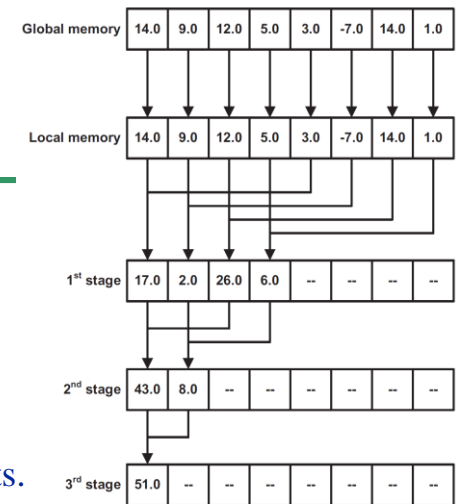


Figure 10.1
Multistage reduction
in OpenCL

```
__kernel void reduction_scalar(__global float* data,  
    __local float* partial_sums, __global float* output) {  
  
    int lid = get_local_id(0);  
    int group_size = get_local_size(0);  
  
    partial_sums[lid] = data[get_global_id(0)];  
    barrier(CLK_LOCAL_MEM_FENCE);  
  
    for(int i = group_size/2; i>0; i >>= 1) {  
        if(lid < i) {  
            partial_sums[lid] += partial_sums[lid + i];  
        }  
        barrier(CLK_LOCAL_MEM_FENCE);  
    }  
  
    if(lid == 0) {  
        output[get_group_id(0)] = partial_sums[0];  
    }  
}
```

From *OpenCL in Action* by M. Scarpino (2012)

1 Read data to
local memory

Perform
reduction stages

Write output to
global memory



• Three reduction methods on the CPU (Serial)

```
void generate_random_float_array(float *array, int n) {
    srand((unsigned int)201803); // Always the same input data
    for (int i = 0; i < n; i++) {
        //array[i] = 2.0f*((float)rand() / RAND_MAX - 0.5f); ← Data A
        array[i] = 100.0f*((float)rand() / RAND_MAX); ← Data B
        //array[i] = 1.0f;
    }
}

float reduction_on_the_CPU(float *array, int n) {
    int i;
    float sum = 0.0f;

    for (i = 0; i < n; i++) {
        sum += array[i];
    }
    return sum;
}

float reduction_on_the_CPU_reduction(float *array, int n) {
    int i, j;
    float sum = 0.0f;

    float *array_b = (float *)malloc(sizeof(float)*n);
    if (array_b == NULL) {
        fprintf(stderr, "+++ Error: cannot allocate memory for array_b...\n");
        exit(EXIT_FAILURE);
    }
    memcpy(array_b, array, sizeof(float)*n);

    for (i = n/2; i > 0; i >>= 1) {
        for (j = 0; j < i; j++) {
            array_b[j] += array_b[j + i];
        }
    }
    sum = array_b[0];
    free(array_b);
    return sum;
}

float reduction_on_the_CPU_KahanSum(float *array, int n) { // From Wikipedia!!!
    int i;
    float sum = 0.0f, c = 0.0f, t, y;

    for (i = 0; i < n; i++) {
        y = array[i] - c;    t = sum + y;
        c = (t - sum) - y;    sum = t;
    }
    return sum;
}
```



• Execution results

[Data A]

The number of elements to be reduced = 16777216

^^^ Test 1: serial reduction on the CPU ^^^

- * Time by host clock for simple reduction = 35.452ms
- * Time by host clock for pairwise reduction = 62.109ms
- * Time by host clock for reduction through Kahan sum = 118.810ms

^^^ Test 2: parallel reduction on the GPU ^^^

KERNEL = reduction_scalar

- * Time by device clock:
 - Time from QUEUED to END = 0.777ms
 - Time from SUBMIT to END = 0.757ms
 - Time from START to END = 0.756ms

+ Check PASSED!
[-3.376091e+03/-3.376028e+03/-3.376028e+03(CPU) = -3.376033e+03(GPU)]

KERNEL = reduction_vector

- * Time by device clock:
 - Time from QUEUED to END = 0.403ms
 - Time from SUBMIT to END = 0.393ms
 - Time from START to END = 0.392ms

+ Check PASSED!
[-3.376091e+03/-3.376028e+03/-3.376028e+03(CPU) = -3.376039e+03(GPU)]

Program ended with exit code: 0

[Data B]

The number of elements to be reduced = 16777216

^^^ Test 1: serial reduction on the CPU ^^^

- * Time by host clock for simple reduction = 35.445ms
- * Time by host clock for pairwise reduction = 59.801ms
- * Time by host clock for reduction through Kahan sum = 119.155ms

^^^ Test 2: parallel reduction on the GPU ^^^

KERNEL = reduction_scalar

- * Time by device clock:
 - Time from QUEUED to END = 0.700ms
 - Time from SUBMIT to END = 0.651ms
 - Time from START to END = 0.650ms

+ Check FAILED!
[8.144400e+08/8.386920e+08/8.386920e+08(CPU) = 8.386916e+08(GPU)]

KERNEL = reduction_vector

- * Time by device clock:
 - Time from QUEUED to END = 0.403ms
 - Time from SUBMIT to END = 0.392ms
 - Time from START to END = 0.391ms

+ Check FAILED!
[8.144400e+08/8.386920e+08/8.386920e+08(CPU) = 8.386927e+08(GPU)]

Program ended with exit code: 0